

Understand Application Security



Elle Krout

Principal Course Author, Pluralsight

Kubernetes Security Contexts



**Security contexts define
privilege and access control
settings for a pod or
container**



Security Context Privilege and Access Control Settings

UID/GID settings

Privilege escalation configuration

Filesystem settings

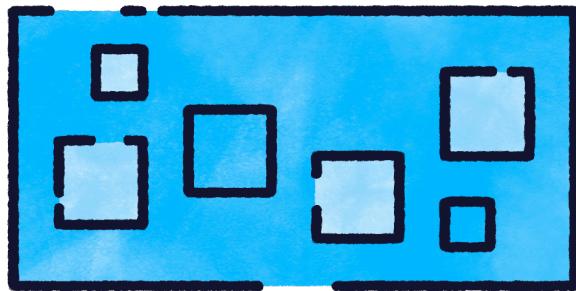
AppArmor/SELinux settings

Seccomp settings

Linux capabilities options



Security Contexts Can Be Set at Two Levels



Pod

Settings apply to all containers in the pod; ideal for settings that needs to be consistent



Container

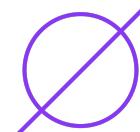
Settings apply to only a single container; ideal for individualized settings



Use Cases



Run as non-root to prevent privilege escalation



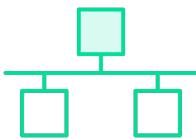
Reduce access to unneeded kernel calls with seccomp



Reduce potential issues by providing read-only access



Limit pod communication with SELinux



Use Linux capabilities to bind containers to a port



... and more!





Security Contexts

Allow you to reduce privileges, enforce consistency, and apply kernel-level protections to your pods and containers



Using Kubernetes Security Contexts



Assign Security Context to Container

secure-container.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  containers:
  - name: app
    image: busybox
    command: ["sh", "-c", "sleep 3600"]
    securityContext:
      allowPrivilegeEscalation: false
      runAsGroup: 3000
      runAsUser: 1000
```



Assign Security Context to Pods

secure-container.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  securityContext:
    runAsGroup: 3000
    runAsUser: 1000
  containers:
    - name: app
      image: busybox
      command: ["sh", "-c", "sleep 3600"]
```



Security Context Parameters

Scope	Matches
<code>runAsUser</code>	Define user ID for container process
<code>runAsGroup</code>	Define group ID for container process
<code>readOnlyRootFilesystem</code>	Force container to have a read-only filesystem
<code>procMount</code>	Define proc mount for container
<code>allowPrivilegeEscalation</code>	Set whether container process can gain more privilege than its parent process
<code>privileged</code>	Boolean to set container to what is essentially root on the host

Security Context Parameters Continued

Scope	Matches
<code>runAsNonRoot</code>	Boolean indicating a container should run as a non-root user
<code>appArmorProfile</code>	Define AppArmor profile and settings
<code>capabilities</code>	Add or drop Linux capabilities
<code>seLinuxOptions</code>	Define SELinux context
<code>seccompProfile</code>	Define seccomp profile and settings
<code>windowsOptions</code>	Security options for Windows containers

View Security Contexts

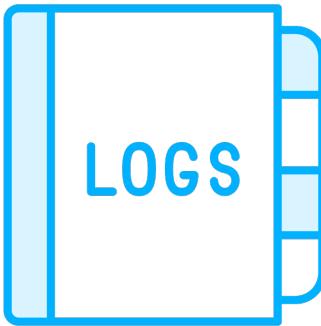
```
> kubectl get pod <pod-name> -o yaml
```

```
...
```

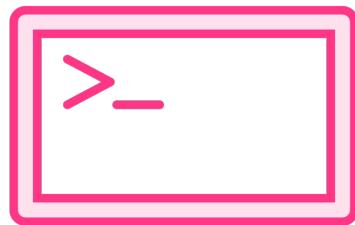
```
spec:  
  containers:  
    - command:  
      - sh  
      - -c  
      - sleep 3600  
    image: busybox  
    imagePullPolicy: Always  
    name: demo  
    resources: {}  
    securityContext:  
      allowPrivilegeEscalation: false  
      runAsGroup: 3000  
      runAsUser: 1000
```



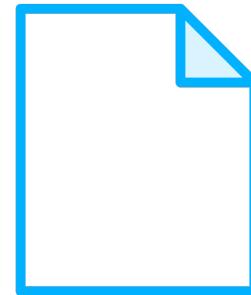
Troubleshoot Security Contexts



Check events and logs



Exec into container

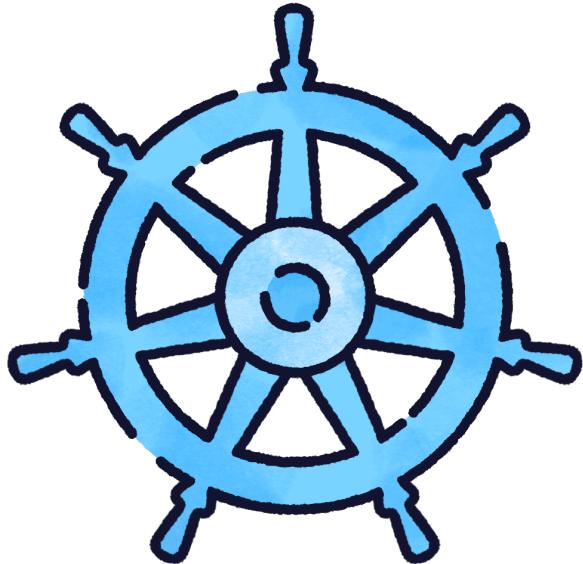


Validate the filesystem



Verify Linux capabilities





Security Contexts...

Are an essential security tool for managing your pod and container privileges and access



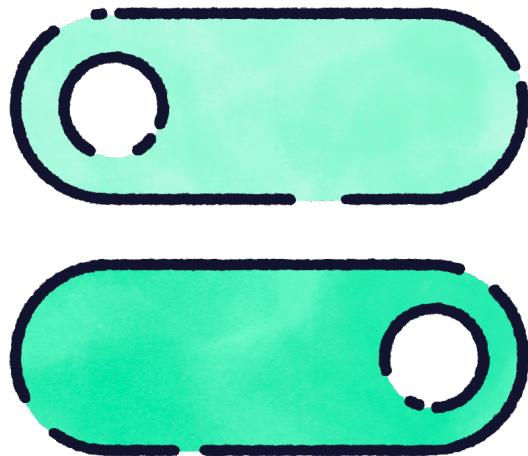
Kubernetes Capabilities



**Linux capabilities are
enhanced features available
on all Linux-based hosts**



Linux Capabilities



A way of breaking down root privileges into smaller permissions

Containers do not start with full privileges by default

Settings determined by runtime



View Capability Options

> `man capabilities`

CAPABILITIES(7)

Linux Programmer's Manual

CAPABILITIES(7)

NAME

`capabilities` - overview of Linux capabilities

DESCRIPTION

For the purpose of performing permission checks, traditional UNIX implementations distinguish two categories of processes: privileged processes (whose effective user ID is 0, referred to as superuser or root), and unprivileged processes (whose effective UID is non-zero). Privileged processes bypass all kernel permission checks, while unprivileged processes are subject to full permission checking based on the process's credentials (usually: effective UID, effective GID, and supplementary group list).



Add or Drop Capabilities

secure-container.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  containers:
  - name: app
    image: busybox
    command: ["sh", "-c", "sleep 3600"]
    securityContext:
      capabilities:
        drop: ["ALL"]
        add: ["NET_BIND_SERVICE"]
```

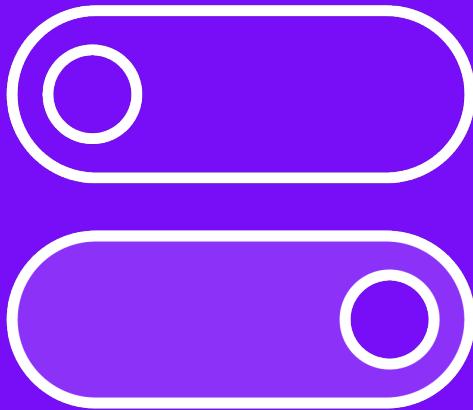


Verify Capabilities are Set

```
> kubectl exec -it <pod> -- sh  
> cat /proc/1/status | grep Cap
```

```
CapInh: 0000000000000000  
CapPrm: 0000000a80425fb  
CapEff: 0000000a80425fb  
CapBnd: 0000000a80425fb  
CapAmb: 0000000000000000
```





Linux Capabilities

Are a powerful way to fine-tune your container
privileges beyond simply letting them run
as root



Demo: Applying Secure Setting in Kubernetes



Exam Scenario



Apply sensible pod-level security default to the monitoring-pod, ensuring it does not allow privilege escalation, and uses the RunTimeDefault seccomp profile

The two containers in the pod, sidecar and logger also need additional security settings. sidecar should run as UID 1000 and drop NET_RAW capabilities. logger should be set to run as both user and group 2000. Both containers should disable privilege escalation



Kubernetes Security Admission



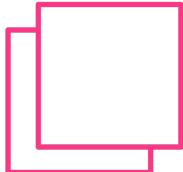
**The PodSecurityAdmission
controller enables security
standards on all pods**



PodSecurityAdmission Permission Levels



Privileged: Most permissive; allows workloads broad access to host



Baseline: Minimally restrictive; blocks known privilege escalations, allows for common container patterns



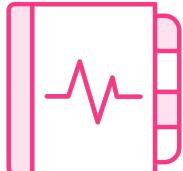
Restricted: Most restrictive; enforces current best practice security suggestions



PodSecurityAdmission Permission Modes



Enforce: Enforces set level; rejects pods in violation



Audit: Records issues in audit log, but lets pods deploy



Warn: Outputs violations to screen but lets pods deploy



Apply Level and Mode to Namespace

```
> kubectl label namespace test pod-security.kubernetes.io/warn=baseline  
  
> kubectl apply -f violation-pod.yaml
```

Warning: would violate PodSecurity "baseline:latest": host namespaces
(hostNetwork=true)



View PodSecurityAdmission Labels for Namespace

```
> kubectl get ns <namespace> --show-labels
```

NAME	STATUS	AGE	LABELS
test	Active	46h	kubernetes.io/metadata.name=test,pod-security.kubernetes.io/warn=baseline



Demo: Working with the Pod Security Admission



Exam Scenario



Create a namespace, `production-test`. This namespace should have the highest level of restriction for its pod security, fully enforced. That said, before this is fully enabled, you should test the `bad-pod.yaml` pod. Have any issues output to the screen after you deployment, as well as tracked in the audit logs

Remove the pod, then apply full enforcement to the namespace. Attempt to deploy the pod again

