

Implement Probes and Health Checks



Elle Krout

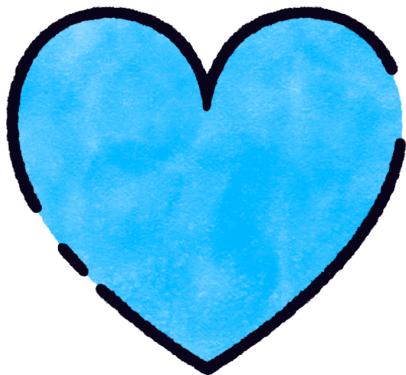
Principal Course Author, Pluralsight

Kubernetes Probes and Health Checks



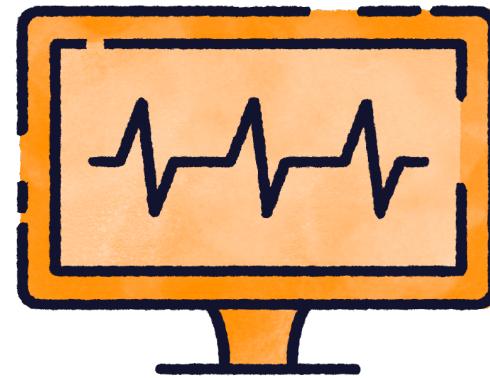
**Observability is the ability
to view and collect data
about your pods and
clusters.**





Health Check

A tool for determining whether an application or service is functioning



Probe

A diagnostic mechanism used by **kubelet** to determine container health



Probes Ensure Pods Are...

Self-healing

Resilient

Stable



Exam Tip!

**Using probes is a key part of passing this section of the CKAD...
not to mention building strong Kubernetes systems in the real
world.**



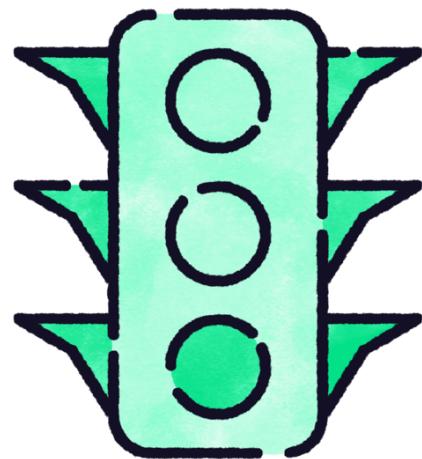
Types of Kubernetes Probes



The Three Probes



Startup



Readiness



Liveness



Startup Probe



Used for prolonged container start times

- Databases
- Java Virtual Machines

Verifies the application on the container has started

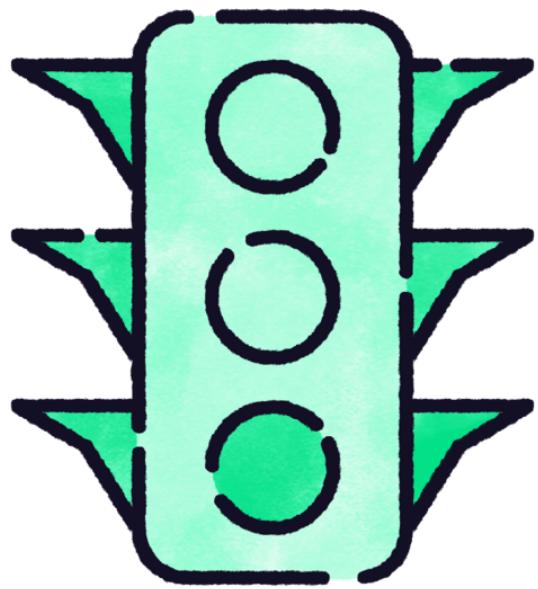
Runs at startup until probe succeeds or fails

Stops running after success/failure

Disables other probes until succeeds



Readiness Probe



Determines if endpoints can receive requests

Runs during containers lifecycle

If fails, removes pod from service

Graceful way of handling:

- Initialization
- Overloaded services



Liveness Probe



Determines if pod is healthy

If fails, destroys pod and creates new one based on restartPolicy

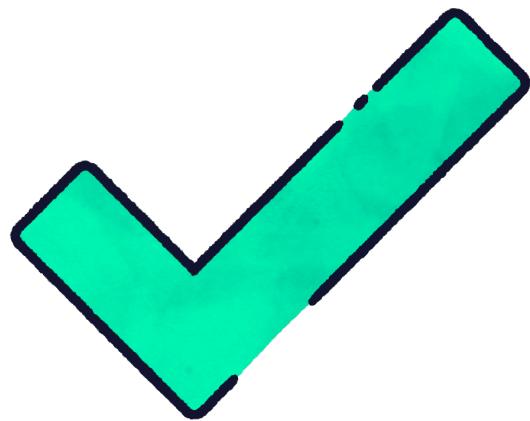
Does not wait for readiness probes to run

Ideal when applications experience:

- Deadlocks
- Hung processes
- Unresponsive services



Probe Results



Success
Everything works!



Failure
Remove or restart

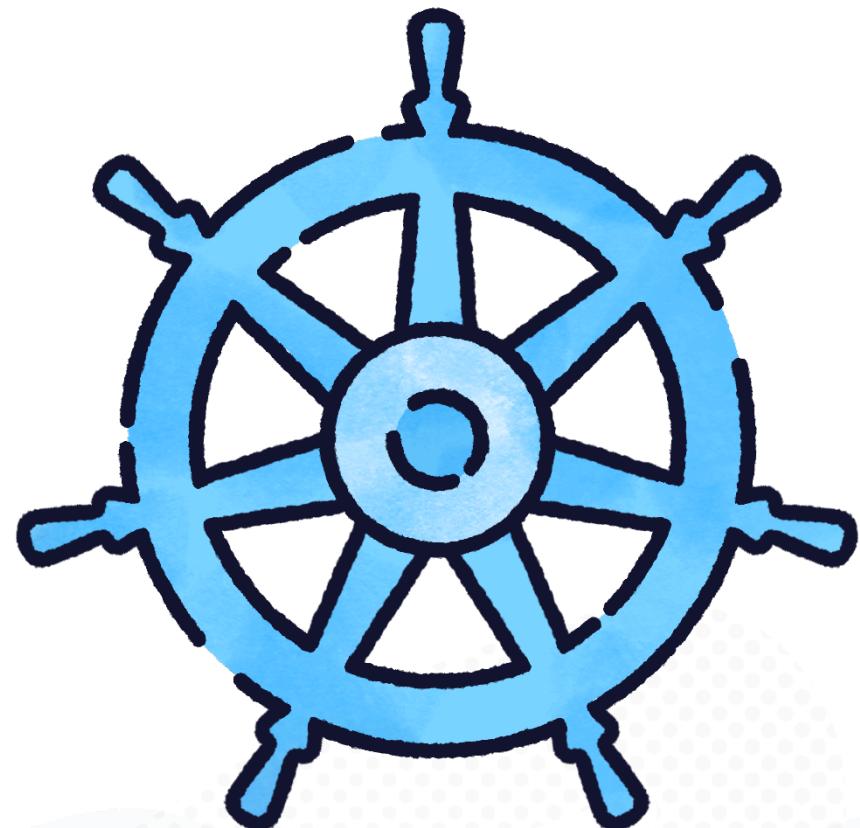


Unknown
Treated as failure



**Probes help build resilient
Kubernetes applications**

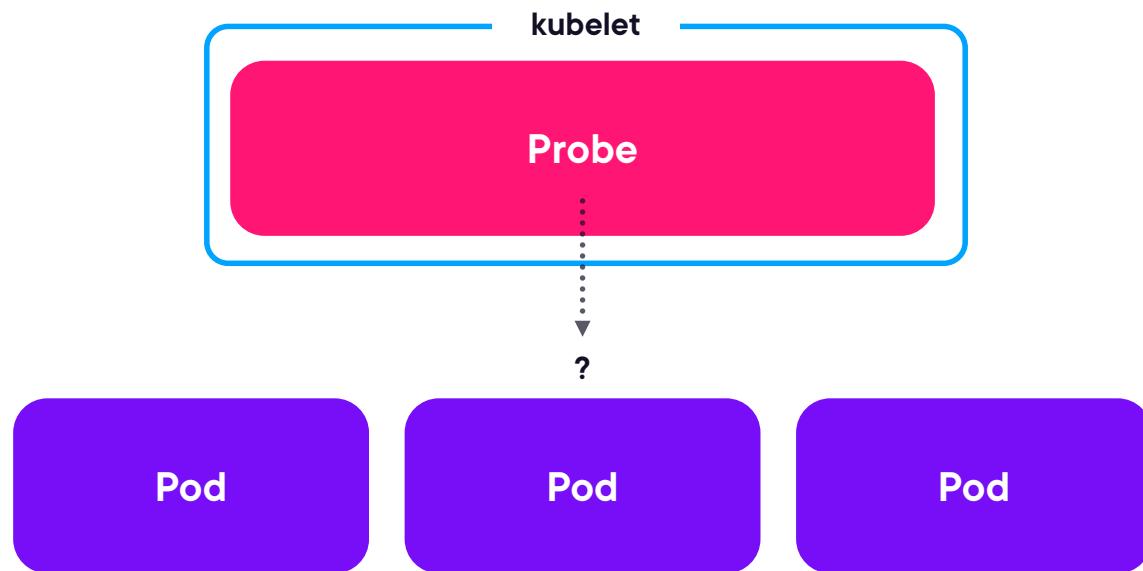
**Know the different and
when to use them to pass
the CKAD exam**



Working with Kubernetes Probes



Kubernetes Probes



httpGet

Probe creates an HTTP GET request against container; expected 200-399 response

pod.yaml

```
livenessProbe:  
  httpGet:  
    port: 80 # Any number between 1-65535; default 80
```



httpGet

Probe creates an HTTP GET request against container; expected 200-399 response

pod.yaml

```
livenessProbe:  
  httpGet:  
    port: 80  
    host: k8s.example.local # Default IP address
```



httpGet

Probe creates an HTTP GET request against container; expected 200-399 response

pod.yaml

```
livenessProbe:  
  httpGet:  
    port: 80  
    host: k8s.example.local  
    scheme: http # Default HTTPS
```



httpGet

Probe creates an HTTP GET request against container; expected 200-399 response

pod.yaml

```
livenessProbe:  
  httpGet:  
    port: 80  
    host: k8s.example.local  
    scheme: http  
    path: /healthz # Default /
```



httpGet

Probe creates an HTTP GET request against container; expected 200-399 response

pod.yaml

```
livenessProbe:  
  httpGet:  
    port: 80  
    host: k8s.example.local  
    scheme: http  
    path: /healthz  
    httpHeaders:  
    - name: Custom  
      value: Hello-Worlds
```



tcpSocket

Checks against a port at the container's IP address

pod.yaml

```
readinessProbe:  
  tcpSocker:  
    port: 8080 # Any number between 1-65535
```



exec

Probe performs action inside of container to check health

pod.yaml

```
startupProbe:  
  exec:  
    command:  
      - cat  
      - /tmp/started
```



Which Handlers Should You Use?



`httpGet` for web apps or application with a health or status endpoint



`tcpSocket` for services with raw TCP endpoints



`exec` for applications where state can only be verified internally



Sample Manifest

```
spec:  
  containers:  
    - name: web-app  
      image: nginx:latest  
      readinessProbe:  
        httpGet:  
          path: /  
          scheme: http  
          port: 80  
      initialDelaySeconds: 2  
      periodSeconds: 5
```



Sample Manifest

```
spec:  
  containers:  
    - name: web-app  
      image: nginx:latest  
      readinessProbe:  
        httpGet:  
          path: /  
          scheme: http  
          port: 80  
      initialDelaySeconds: 2  
      periodSeconds: 5
```



Sample Manifest

```
spec:  
  containers:  
    - name: web-app  
      image: nginx:latest  
      readinessProbe:  
        httpGet:  
          path: /  
          scheme: http  
          port: 80  
      initialDelaySeconds: 2 # How long in seconds before probe runs  
      periodSeconds: 5
```



Sample Manifest

```
spec:  
  containers:  
    - name: web-app  
      image: nginx:latest  
      readinessProbe:  
        httpGet:  
          path: /  
          scheme: http  
          port: 80  
      initialDelaySeconds: 2  
      periodSeconds: 5 # How often the probe will run
```



Sample Manifest

```
spec:  
  containers:  
    - name: web-app  
      image: nginx:latest  
      readinessProbe:  
        httpGet:  
          path: /  
          scheme: http  
          port: 80  
      initialDelaySeconds: 2  
      periodSeconds: 5  
      timeoutSeconds: 5 # How long the probe will wait before timeout
```



Sample Manifest

```
spec:  
  containers:  
    - name: web-app  
      image: nginx:latest  
      readinessProbe:  
        httpGet:  
          path: /  
          scheme: http  
          port: 80  
      initialDelaySeconds: 2  
      periodSeconds: 5  
      timeoutSeconds: 5  
      failureThreshold: 2 # How many times a probe can fail
```



Sample Manifest

```
spec:  
  containers:  
    - name: web-app  
      image: nginx:latest  
      readinessProbe:  
        httpGet:  
          path: /  
          scheme: http  
          port: 80  
      initialDelaySeconds: 2  
      periodSeconds: 5  
      timeoutSeconds: 5  
      failureThreshold: 2  
      successThreshold: 2 # How many times a probe needs to succeed for success
```



Sample Manifest

```
spec:
```

```
  containers:
```

```
    - name: web-app
```

```
      image: nginx:latest
```

```
      readinessProbe:
```

```
        httpGet:
```

```
          path: /
```

```
          scheme: http
```

```
          port: 80
```

```
        initialDelaySeconds: 2
```

```
        periodSeconds: 5
```

```
        timeoutSeconds: 5
```

```
        failureThreshold: 2
```

```
        successThreshold: 2
```

```
      terminationGracePeriodSeconds: 20 # Time between requested shutdown and actual  
                                         shutdown
```



Demo: Adding and Testing Kubernetes Probes



Exam Scenario



The `slow-start.yaml` pod manifest contains the configuration for a slow starting application that listens on port 9090. The connection starts before the service is fully ready; a file under `/tmp/live` indicates the pod has fully started up.

Add a combination of startup, readiness, and liveness probes to the manifest.

Ensure the startup probe confirms the pod is ready by accessing the `/tmp/live` and wait 30 seconds before it is run. The startup probe should be able to fail up to 5 times.

The readiness probe should succeed twice for the pod to be considered ready. It should check the connection at 9090.

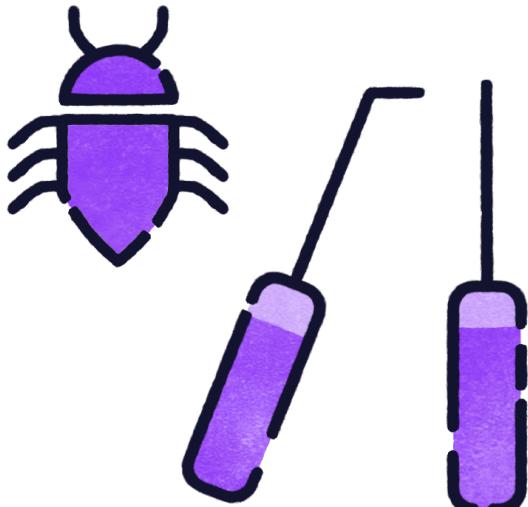
Delay the liveness probe for 10 seconds. Have it run every 30 seconds after that. It should similarly check the connection at 9090.



Troubleshooting Kubernetes Probes



Troubleshooting Tips



Check the probe configuration

Check events:

```
kubectl describe pod <pod-name>
```

Check cluster events:

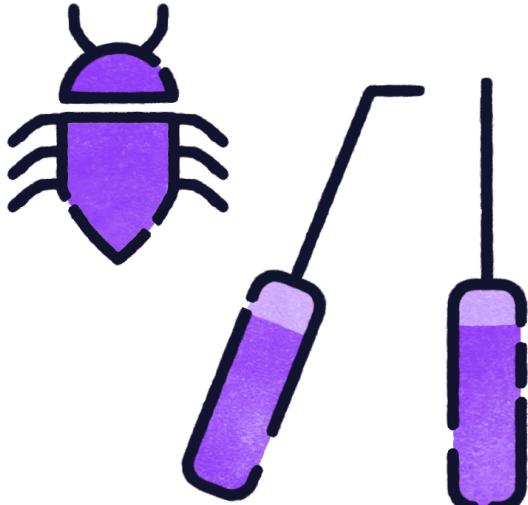
```
kubectl get events --sort-by= '.lastTimestamp'
```

Check logs:

```
kubectl logs <pod-name> [-c <container-name>]
```



Troubleshooting Tips



Check the probe configuration

Check events:

```
kubectl describe pod <pod-name>
```

Check cluster events:

```
kubectl get events --sort-by='.lastTimestamp'
```

Check logs:

```
kubectl logs <pod-name> [-c <container-name>] --previous
```

