

Use Built-in Tools to Monitor Kubernetes Applications



Elle Krout

Principal Course Author, Pluralsight

Monitoring in Kubernetes



Monitoring Benefits

Operational
Insights

Performance
Tuning

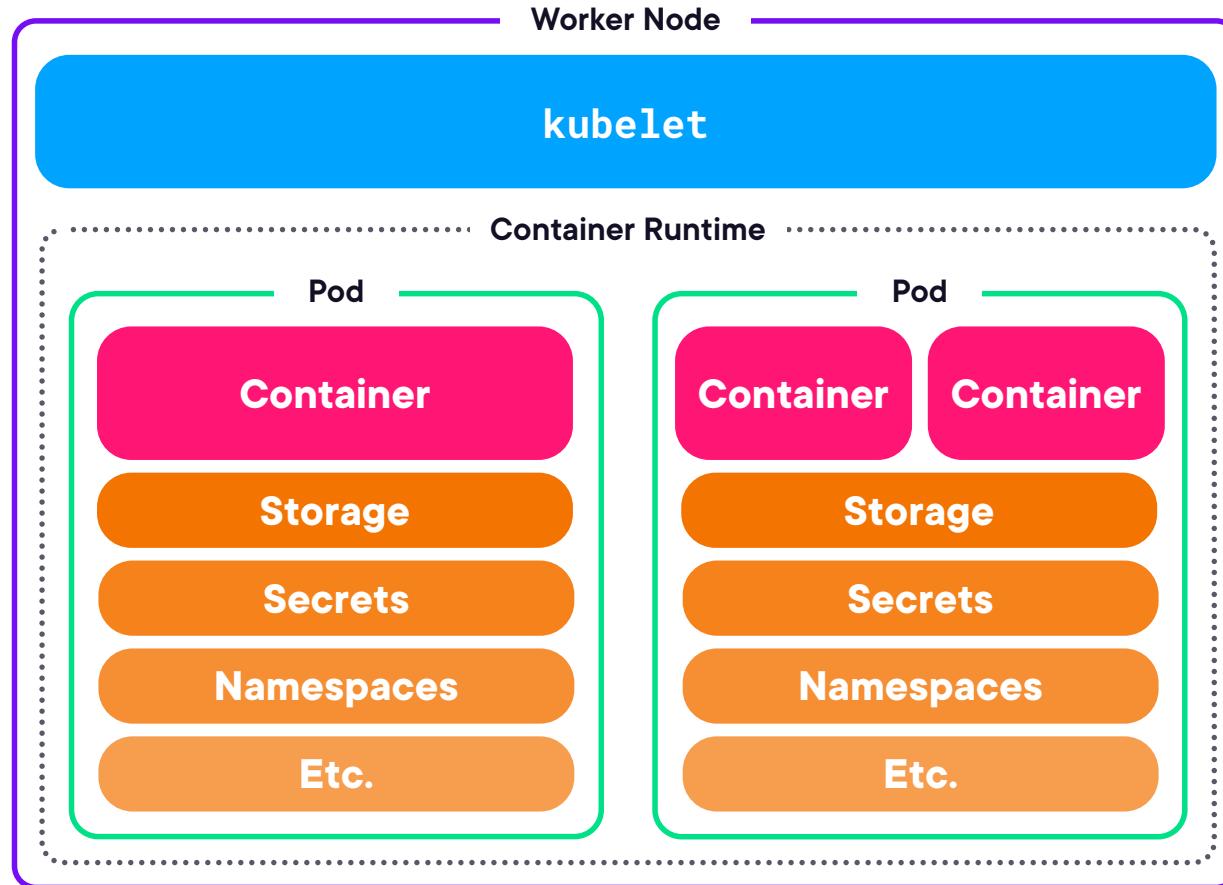
Troubleshooting
Assistance

Scaling Support

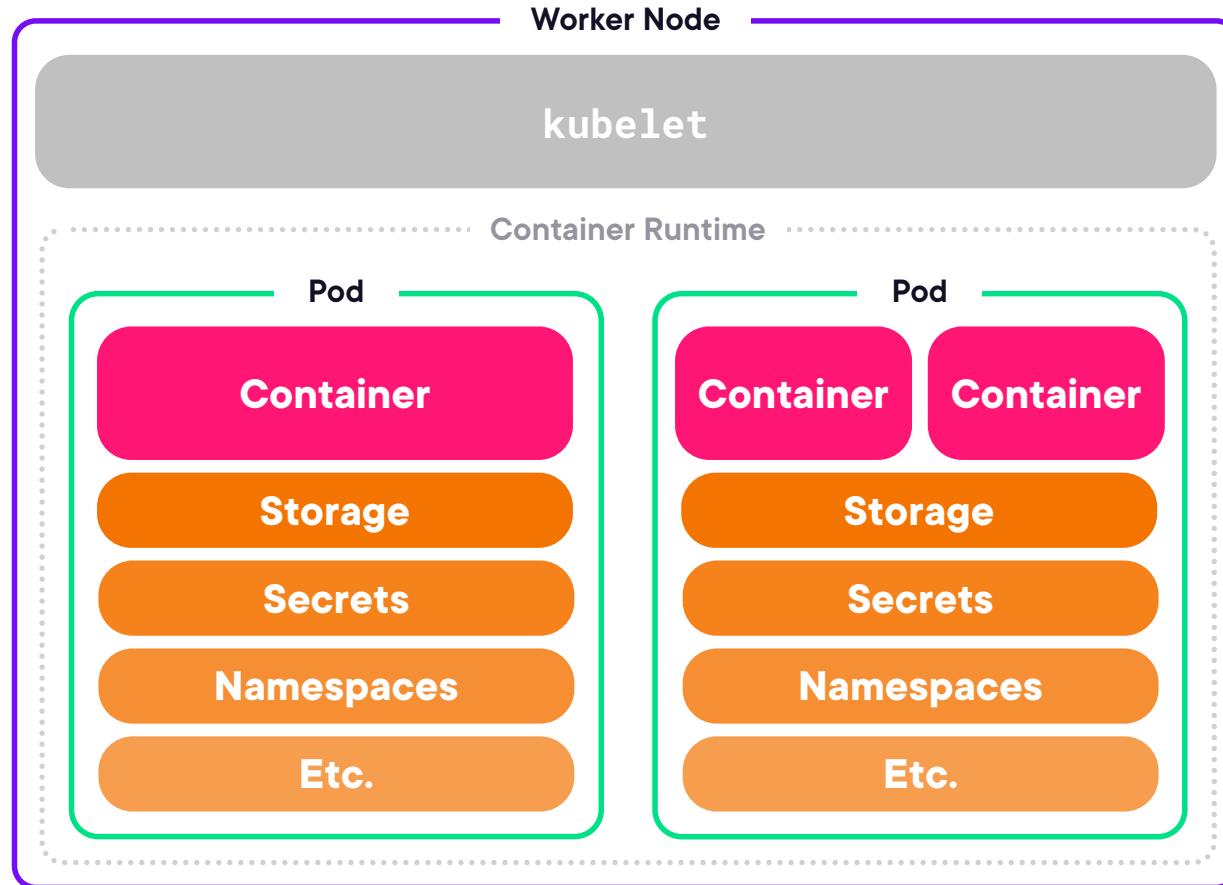
Data-driven
Decisions



What Should You Monitor?



What Should You Monitor?



CLI-only monitoring

No outside monitoring solutions

Metrics Server

Metrics API



Metrics Server is a Kubernetes-provided in-cluster data aggregator that collects resource metrics and exposes them to the API server through the use of the Metrics API.



**Use /metrics endpoint for
sending metrics to other
services.**



Metrics Server

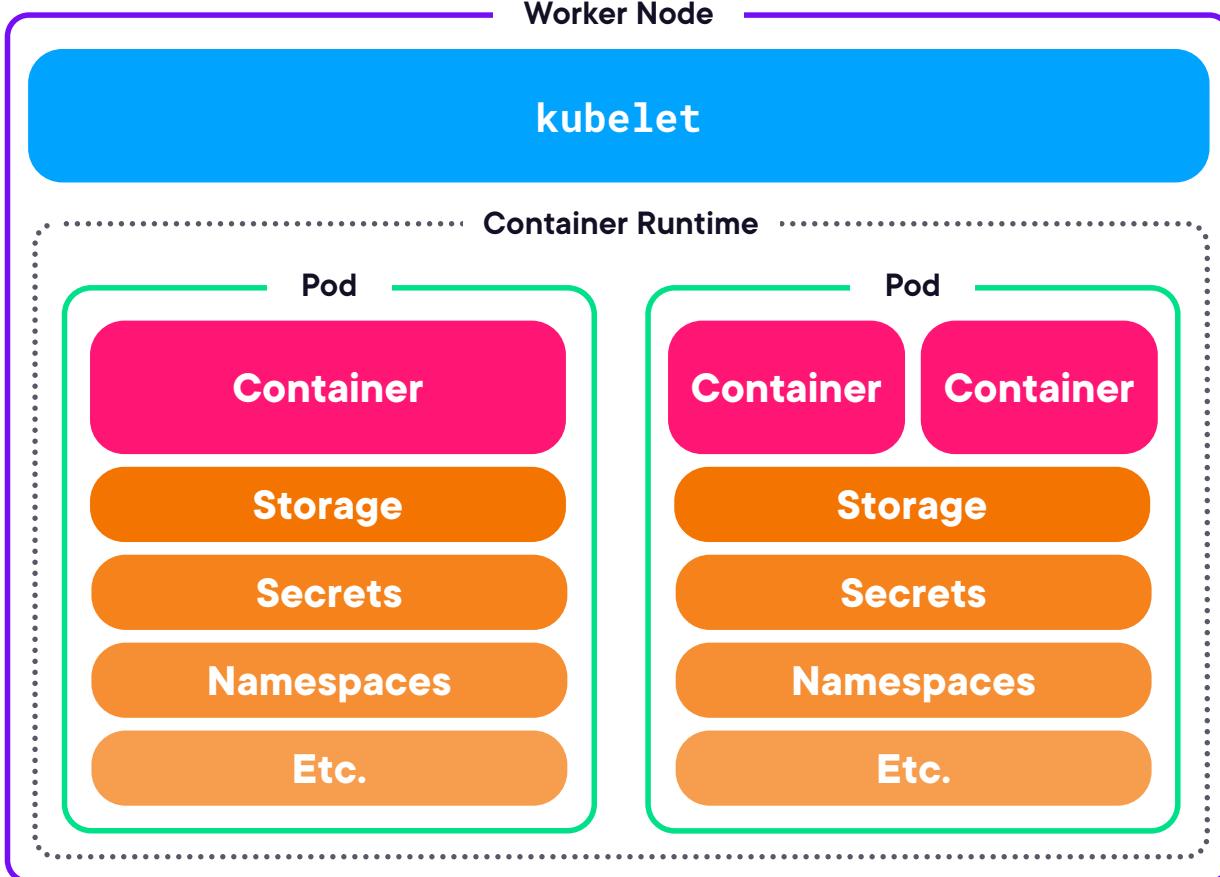


Not installed by default, but will be on the exam clusters

View resource-level metrics

- CPU
- Memory





Demo: Installing Metrics Server



Metrics Server Prerequisites

The kube-apiserver must enable an aggregation layer



Metrics Server Prerequisites

**Nodes must have Webhook authentication and
authorization enabled**



Metrics Server Prerequisites

The kubelet certificate needs to be signed by cluster Certificate Authority; alternatively, you can disable certificate validation by passing adding `--kubelet-insecure-tls` to the Metrics Server startup command in the manifest



Metrics Server Prerequisites

The container runtime must implement a container metrics RPCs (or have cAdvisor support)



Metrics Server Prerequisites

The network should support allow the control plane node to reach the Metrics Server on port 10250

Metrics server also needs to be able to speak to kubelet on all nodes



Accessing Kubernetes Metrics Data



kubectl top



View Metrics Server Data

```
> kubectl top nodes
```

NAME	CPU(cores)	CPU(%)	MEMORY(bytes)	MEMORY(%)
controller	217m	10%	1638Mi	43%
worker1	46m	2%	896Mi	23%



View Metrics Server Data

```
> kubectl top pods
```

NAME	CPU(cores)	MEMORY(bytes)
web-app-7b99994d65-lwcmx	1m	5Mi
web-app-7b99994d65-st9rx	1m	9Mi
web-app-7b99994d65-zhbqd	1m	12Mi

Metrics Server does not provide historical data



View Metrics Server Data

```
> kubectl top pods -n kube-system
```

NAME	CPU(cores)	MEMORY(bytes)
coredns-674b8bbfcf-kwpkd	3m	41Mi
coredns-674b8bbfcf-nbr4w	2m	41Mi
etcd-controller	26m	79Mi
kube-apiserver-controller	64m	272Mi
kube-controller-manager-controller	25m	124Mi
kube-proxy-rhz2v	1m	77Mi
kube-proxy-sr4sw	1m	76Mi
kube-scheduler-controller	12m	80Mi
metrics-server-5b5786cf87-4p4wj	5m	66Mi



View Metrics Server Data

```
> kubectl top pods --all-namespaces
```

NAMESPACE	NAME	CPU(cores)	MEMORY(bytes)
default	web-app-7b99994d65-lwcmx	1m	5Mi
default	web-app-7b99994d65-st9rx	1m	9Mi
default	web-app-7b99994d65-zhbqd	1m	12Mi
kube-flannel	kube-flannel-ds-6bddn	14m	54Mi
kube-flannel	kube-flannel-ds-9tzzq	18m	54Mi
kube-system	coredns-674b8bbfcf-kwpkd	2m	41Mi
kube-system	coredns-674b8bbfcf-nbr4w	3m	41Mi
kube-system	etcd-controller	29m	80Mi
kube-system	kube-apiserver-controller	51m	272Mi
kube-system	kube-controller-manager-controller	26m	124Mi
kube-system	kube-proxy-rhz2v	1m	77Mi
...			



View Raw Metrics Data

```
> kubectl get --raw /api/v1/nodes/<NODE_NAME>/proxy/metrics/resource

# HELP container_cpu_usage_seconds_total [STABLE] Cumulative cpu time consumed by the
container in core-seconds
# TYPE container_cpu_usage_seconds_total counter
container_cpu_usage_seconds_total{container="kube-flannel",namespace="kube-
flannel",pod="kube-flannel-ds-6bddn"} 8.782792 1750962620660
container_cpu_usage_seconds_total{container="kube-proxy",namespace="kube-system",pod="kube-
proxy-sr4sw"} 1.187455 1750962623858
container_cpu_usage_seconds_total{container="nginx",namespace="default",pod="web-
app-7b99994d65-lwcmx"} 0.151779 1750962627564
container_cpu_usage_seconds_total{container="nginx",namespace="default",pod="web-
app-7b99994d65-st9rx"} 0.187888 1750962622829
# HELP container_memory_working_set_bytes [STABLE] Current working set of the container in
bytes
# TYPE container_memory_working_set_bytes gauge
...
```



Demo: Working with Kubernetes Metrics Server



Exam Scenario



Create a web application deployment using the `nginx` image; name the deployment `nginx`, create five replicas, and expose port 80

Check the resource usage for the deployment pods and the related worker nodes

Generate some traffic using `curl` against the `nginx` deployment

Review the resource usage once more

View the raw metrics data for one of the affected nodes. Save it in the file `workernode.txt`



Kubernetes Rollout Monitoring



View Live Rollout Data

```
> kubectl rollout status deployment <DEPLOYMENT_NAME>
```

```
kubectl rollout status deployment <DEPLOYMENT_NAME>
```

```
Waiting for deployment "deployment-name" rollout to finish: 0 of 3 updated replicas  
are available...
```

```
Waiting for deployment "deployment-name" rollout to finish: 1 of 3 updated replicas  
are available...
```

```
Waiting for deployment "deployment-name" rollout to finish: 2 of 3 updated replicas  
are available...
```

```
deployment "deployment-name" successfully rolled out
```



View Historical Rollouts

```
> kubectl rollout history deployment <DEPLOYMENT_NAME>
```

```
REVISION  CHANGE-CAUSE
```

```
1        <none>
2        kubectl set image deployment <deployment-name> busybox=busybox:latest --
record=true
```



View Historical Rollouts

```
> kubectl rollout history deployment <DEPLOYMENT_NAME>
```

```
REVISION  CHANGE-CAUSE
```

```
1        <none>
```

```
2        kubectl set image deployment <deployment-name> busybox=busybox:latest --  
record=true
```



Demo: Monitoring Kubernetes Rollouts



**[https://github.com/
pluralsight-cloud/
application-observability-
maintenance-ckad](https://github.com/pluralsight-cloud/application-observability-maintenance-ckad)**



Exam Scenario



Deploy the web-app manifest and view the rollout deployment live

Upgrade the deploy to use the latest version of the Nginx image. Watch the deployment, then confirm by viewing the deployment history

